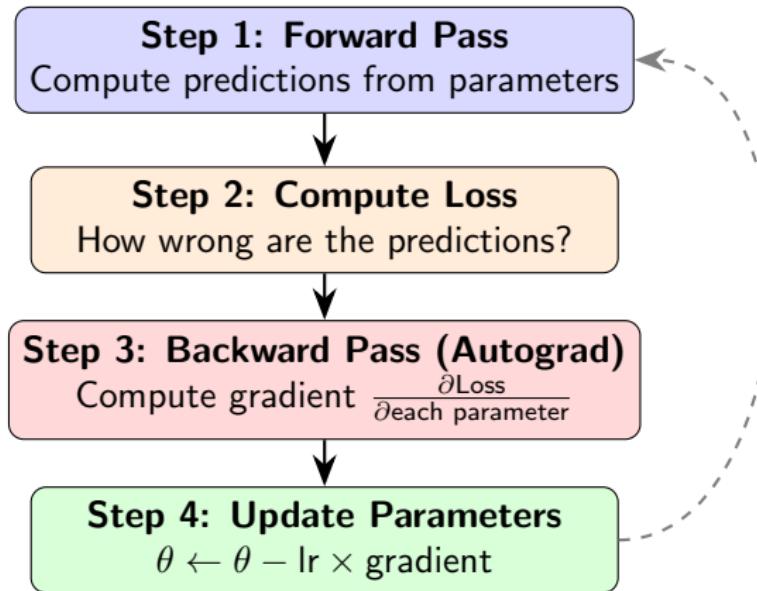


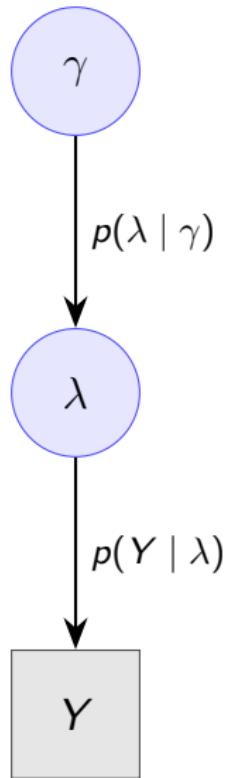
## Why $\gamma$ Doesn't Learn: Centered vs. Non-Centered MAP

# How ML Training Works (4 Steps, Every Epoch)



**Key rule:** In Step 3, a parameter only gets a gradient from the loss if it was **used in Step 1** to compute the prediction. If a parameter isn't in the forward pass, autograd can't trace back to it.

# The Hierarchical Model



**ALADYN:**

►  $\gamma$  ≡ genetic effects on disease signatures (hyperparameter)

# Centered Model: The Forward Pass

**What the computer does each epoch:**

1. Read  $\lambda$  from memory (it's a free `nn.Parameter`)

2.  $\theta = \text{softmax}(\lambda)$  (mixing proportions)

3.  $\pi = \theta \cdot \phi \cdot \kappa$  (disease probabilities)

4. Loss =  $\text{NLL}(Y, \pi) + w \cdot (\lambda - \mu(\gamma))^{\top} \Omega^{-1} (\lambda - \mu(\gamma))$

**Notice:**  $\gamma$  appears **only in the prior term** (Step 4), never in the forward computation (Steps 1–3).

$\lambda$  is read from memory  $\rightarrow \theta \rightarrow \pi \rightarrow \text{NLL}$ .

The chain  $Y \rightarrow \pi \rightarrow \theta \rightarrow \lambda$  does not pass through  $\gamma$ .

# Reparameterized Model: The Forward Pass

**What the computer does each epoch:**

1. Read  $\delta$  and  $\gamma$  from memory (both are `nn.Parameters`)

2.  $\lambda = G \cdot \gamma + \delta$  ( $\gamma$  enters the forward pass!)

3.  $\theta = \text{softmax}(\lambda)$  (mixing proportions)

4.  $\pi = \theta \cdot \phi \cdot \kappa$  (disease probabilities)

5. Loss =  $\text{NLL}(Y, \pi) + w \cdot \delta^\top \Omega^{-1} \delta$

**Now:**  $\gamma$  is in the forward computation (Step 2).

The chain is:  $\gamma \rightarrow \lambda \rightarrow \theta \rightarrow \pi \rightarrow \text{NLL}$ .

Autograd traces back through this chain, so  $\gamma$  gets a gradient from the data.

# The Backward Pass: Why $\gamma$ Gets Left Out

**Centered model** — backward pass computes:

$$\frac{\partial \text{Loss}}{\partial \lambda} = \underbrace{\frac{\partial \text{NLL}}{\partial \lambda}}_{\text{data signal (strong)}} + \underbrace{w \cdot \frac{\partial \text{prior}}{\partial \lambda}}_{\text{prior signal (weak)}}$$

$$\frac{\partial \text{Loss}}{\partial \gamma} = \underbrace{\frac{\partial \text{NLL}}{\partial \gamma}}_{\text{Not in forward pass!}} + \underbrace{w \cdot \frac{\partial \text{prior}}{\partial \gamma}}_{\text{prior only } (w=1e-4)}$$

# The Backward Pass: Why $\gamma$ Gets Left Out

**Centered model** — backward pass computes:

$$\frac{\partial \text{Loss}}{\partial \lambda} = \underbrace{\frac{\partial \text{NLL}}{\partial \lambda}}_{\text{data signal (strong)}} + \underbrace{w \cdot \frac{\partial \text{prior}}{\partial \lambda}}_{\text{prior signal (weak)}}$$

$$\frac{\partial \text{Loss}}{\partial \gamma} = \underbrace{\frac{\partial \text{NLL}}{\partial \gamma}}_{\text{Not in forward pass!}} + \underbrace{w \cdot \frac{\partial \text{prior}}{\partial \gamma}}_{\text{prior only } (w=1e-4)}$$

---

**Reparam model** — backward pass computes:

$$\frac{\partial \text{Loss}}{\partial \gamma} = \underbrace{\frac{\partial \text{NLL}}{\partial \pi} \cdot \frac{\partial \pi}{\partial \theta} \cdot \frac{\partial \theta}{\partial \lambda} \cdot \frac{\partial \lambda}{\partial \gamma}}_{\text{chain rule through forward pass!}} + 0$$

$\gamma$  is in the forward pass, so the NLL gradient flows back to it.

## Same Objective, Different Computation

Both models minimize the **same function**:

$$p(Y | \lambda) \cdot p(\lambda | \gamma)$$

# Same Objective, Different Computation

Both models minimize the **same function**:

$$p(Y | \lambda) \cdot p(\lambda | \gamma)$$

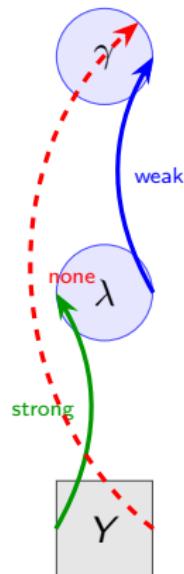
The difference is purely in **how the computer represents  $\lambda$** :

	Centered	Non-centered (Reparam)
Free parameters	$\lambda, \gamma$	$\delta, \gamma$
Forward pass	$\theta = \text{softmax}(\lambda)$	$\lambda = G\gamma + \delta$ $\theta = \text{softmax}(\lambda)$
NLL	$-\log p(Y   \lambda)$	$-\log p(Y   G\gamma + \delta)$
Prior	$(\lambda - G\gamma)^T \Omega^{-1} (\lambda - G\gamma) = \delta^T \Omega^{-1} \delta$	(identical)
$\gamma$ in NLL?	No	Yes
$\gamma$ gets data grad?	No	Yes

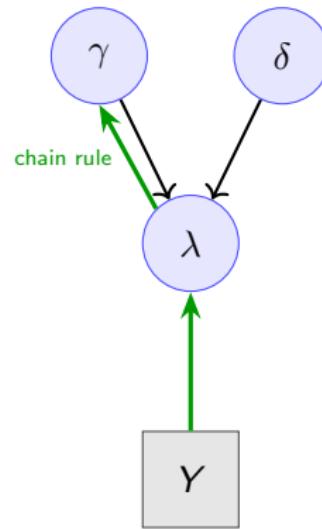
**At any solution:**  $\delta^* = \lambda^* - G\gamma^*$ . If the problem were convex, both would find the same optimum. Because it is **non-convex**, different gradient flows  $\rightarrow$  different local optima.

# The Gradient Flow Picture

Centered:



Non-centered:



Both  $\gamma$  and  $\delta$  feed into  $\lambda$  in the forward pass  $\rightarrow$  both get NLL gradients.

In centered,  $\lambda$  is a dead-end: gradients stop there and never reach  $\gamma$  through the NLL.

# A Simple Analogy

Minimize  $f(x) = (x - 3)^2$  with a preference that  $x$  be near  $\mu$ :

## Centered:

- ▶ Parameters:  $x$  and  $\mu$
- ▶ Loss =  $(x - 3)^2 + w \cdot (x - \mu)^2$
- ▶  $\mu$  only appears in the penalty. If  $w$  is small,  $\mu$  barely moves.

## Non-centered:

- ▶ Parameters:  $z$  and  $\mu$ , where  $x = \mu + z$
- ▶ Loss =  $(\mu + z - 3)^2 + w \cdot z^2$
- ▶  $\mu$  appears in the **main loss**. It gets a strong gradient.

Same answer ( $x^* = 3$ ), but in the second form  $\mu$  learns where  $x$  wants to be. In the first form,  $\mu$  barely moves because  $w$  is tiny.

# Why It Usually Doesn't Matter (for Prediction) $\lambda$ compensates.

For prediction, we:

1. Fix all population parameters  $(\phi, \gamma, \psi, \kappa)$  from training
2. For each new individual: re-estimate  $\lambda_{\text{new}}$  (200 epochs)
3.  $\lambda_{\text{new}}$  has  $K \times T$  free parameters — enough to absorb the signal

So even if  $\gamma \approx 0$  (centered),  $\lambda_{\text{new}}$  can still fit the data well.

# Why It Usually Doesn't Matter (for Prediction)

$\lambda$  compensates.

For prediction, we:

1. Fix all population parameters  $(\phi, \gamma, \psi, \kappa)$  from training
2. For each new individual: re-estimate  $\lambda_{\text{new}}$  (200 epochs)
3.  $\lambda_{\text{new}}$  has  $K \times T$  free parameters — enough to absorb the signal

So even if  $\gamma \approx 0$  (centered),  $\lambda_{\text{new}}$  can still fit the data well.

**However**, it does still matter because:

- Reparam initializes  $\lambda = G\gamma + \delta$  (starting near the genetic prediction)
- Centered initializes  $\lambda$  near zero (no genetic information at start)
- Different starting points  $\rightarrow$  different optimization paths  $\rightarrow$  different final  $\lambda$
- Different trained  $\phi/\psi/\kappa$  between the two pipelines also contribute

# Prediction AUC: Centered vs. Non-Centered

50,000 held-out patients, 100 bootstrap replicates, 28 diseases:

Metric	Centered	Non-centered	$\Delta$	Wins
Static 10-year	0.622	0.654	+0.032	25/28
Dynamic 10-year	0.624	0.629	+0.005	18/28
Dynamic 1-year	0.765	0.883	+0.118	24/28

**Non-centered wins on all 3 metrics.**

- ▶ Largest advantage at short horizon (1-year: +0.118)
- ▶ Smallest advantage at dynamic 10-year (+0.005)
- ▶ **Caveat:** confounded by different population parameters ( $\phi, \psi, \kappa$ ) between the two training pipelines — not purely a parameterization effect

# Trained Parameter Differences

	Centered	Non-centered
Mean $ \gamma $	0.006	0.081
$\kappa$	2.93	4.52
$\gamma$ correlation	0.37 (very different)	
$\psi$ correlation	0.76 (moderately different)	
$\phi$ correlation	0.94 (similar)	

- ▶  $\phi$  (disease trajectories): nearly identical — both fit the data well
- ▶  $\psi$  (signature assignments): moderately different — non-centered less stable
- ▶  $\gamma$  (genetic effects): **centered**  $\gamma \approx 0$  ( $14\times$  smaller)
- ▶ The **prediction AUC difference is driven by all of these** (especially  $\phi/\psi/\kappa$ ), not just  $\gamma$

## Summary

1. **ML training = forward → loss → backward → update.** A parameter only gets a data gradient if it was in the forward pass.
2. Both models optimize the same posterior:  $p(Y|\lambda) \cdot p(\lambda|\gamma)$ .  
The difference is whether you write  $\lambda = G\gamma + \delta$  in the forward step.
3. **Centered** (original):  $\lambda$  is free,  $\gamma$  only in the prior  $\rightarrow \gamma$  gets no data gradient.
4. **Non-centered** (reparam):  $\lambda = G\gamma + \delta$ ,  $\gamma$  in forward pass  $\rightarrow \gamma$  gets data gradient via chain rule.
5. Both are standard approaches. Centered is the default in Stan/Ime4/PyMC. Non-centered is the standard alternative.
6. For prediction: non-centered wins on AUC (all 3 metrics), but confounded by different population parameters.